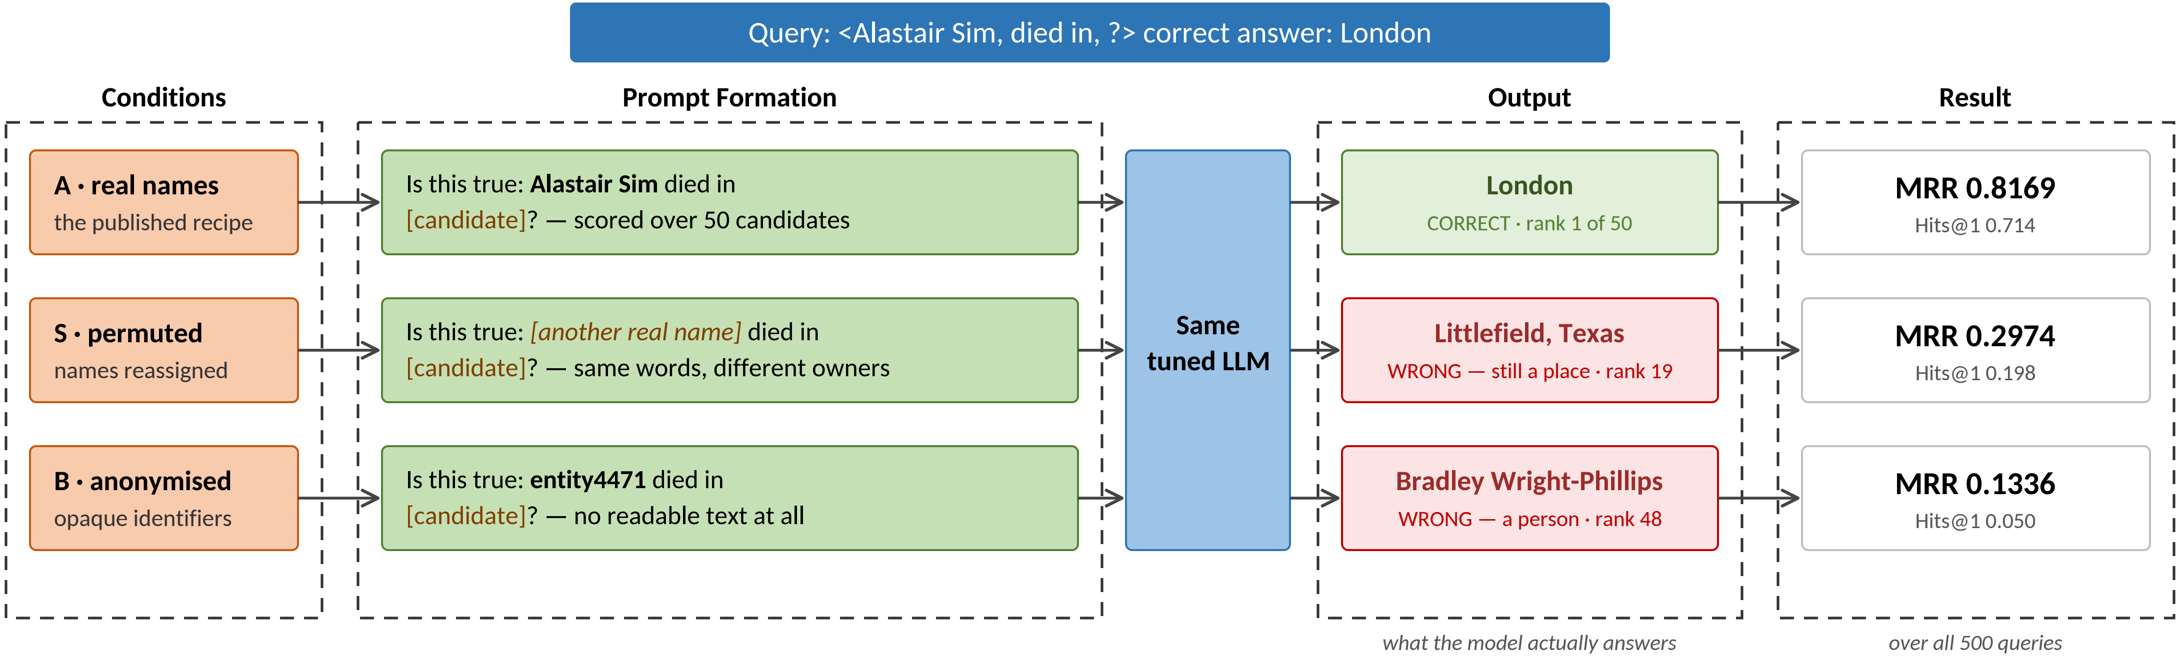

Decomposition of above-chance skill (50-way filtered, chance MRR = 0.0900)

The three arms differ only in entity surface form: relations, graph structure, adaptation recipe and token budget are identical. The **gap between arms** is the measurement, not any single accuracy. Under S the model still returns a **place** — the type survives, the identity does not; under B it returns a person, so readability went too.

Memorisation share = $(\mathrm{MRR}_A - \mathrm{MRR}_B) / (\mathrm{MRR}_A - \mathrm{chance})$ = **94.0 %** on YAGO3-10. [another real name] = a different entity's real name, under a fixed derangement.